

Implementing compact tree-hashing commitment verification in ZK-SNARK

Stefano Trevisani

May 31, 2022

Contents

1	Introduction	3
2	Preliminaries	4
2.1	Computational models and complexity	4
2.2	Finite fields and arithmetic circuits	4
2.3	Rank-1 Constraint systems	6
2.4	Quadratic Arithmetic Programs	8
2.5	Hash functions	10
2.6	Tree hash modes	12
3	ZK-SNARK	14
3.1	Zero Knowledge Proofs	14
3.2	The Pinocchio Protocol	15
4	The libsnark library	17
5	Experiments	17
6	Conclusions and future directions	17

1 Introduction

One of the biggest revolutions of the last decade has been the widespread adoption of blockchain-based technologies. In particular, cryptocurrencies like Bitcoin or Ethereum are widely known even among the non-crypto-enthusiasts and a huge market is growing around them. There are highly interesting applications of the blockchain even outside the financial world: in fact, anything which requires some degree of ‘verifiability’ in a non-trusted environment can benefit from using a blockchain. Typically, a block of the chain does not correspond to a single transaction, as the blockchain would become too big to be stored: many transactions are inserted into a tree data-structure, called Merkle Tree (MT), which is computed bottom-up and whose root is then inserted into the blockchain. In a traditional setup (e.g. the aforementioned cryptocurrencies), all the data which is required to build a block of the blockchain is of public knowledge: when an Ethereum transaction happens, the details of that transaction are shared with the network for it to be validated. Of course, there are many scenarios where one would like the data to remain secret, as there could be risks for privacy (e.g. transactions) or for intellectual property (e.g. algorithms).

Zero-Knowledge Succinct Non-interactive ARgument of Knowledge (ZK-SNARK) systems are cryptographic frameworks which allow for a party to convince other parties that he ‘knows something’ without revealing anything else. For example, one could convince other users that the hash of a transaction is valid without revealing the details of that transaction. ZK systems work over prime fields, and hashing algorithms like SHA-256, which are quite efficient in a ‘vanilla’ scenario, can become extremely slow when translated. For this reason, new hashing algorithms like MiMCHash have been designed with the ZK scenario in mind, and Augmented Binary tRee (ABR) tries to improve Merkle Trees by processing more transactions without requiring more hash function calls.

In the remainder of the report, for clarity we will always have in mind the cryptocurrency scenario, but we want to make clear that all the other use-cases are the same¹. In Section 2, we introduce hash functions and tree-based modes of hashing. In Section 3, we introduce ZK-SNARK systems and briefly describe how they can be implemented. In Section 4, we introduce `libsark`, a C++ library that provides many facilities to implement ZKSNARK algorithms with relative ease, and we will discuss the implementation of MiMC, MTs and ABRs using this library. In Section 5, we perform some experiments to compare MiMC with SHA, and MTs with ABRs in a ZK setting. Finally, in Section 6 we draw our conclusions and discuss possible future work directions.

¹Up to isomorphism, of course.

2 Preliminaries

In this section we are going to introduce some fundamental concepts; while some are relatively basic and wide-known, it can still be useful to skim over them to be sure of having a firm grasp on the main ideas behind ZK-SNARK systems.

2.1 Computational models and complexity

A *computational model* (or model of computation) is any kind of ‘device’, either physical or mathematical, which is able to compute algorithms to solve problems. A particularly interesting class of problems are *decision problems*, the ones that can be answered with ‘yes’ (or ‘accept’, or \top) or ‘no’ (or ‘reject’, or \perp). Every computational model is able to *decide* only a subclass of all decision problems, and even then, not all can be solved *efficiently*, that is, by using an amount of resources (typically, time and space) which is upper-bounded by some polynomial function of the input length. Problems for which a polynomial bound doesn’t exist or isn’t known are said to be *hard* for the computational model. For example, finding solutions to boolean equations (the SAT problem) is believed to be hard for deterministic Turing machines, but it is easy for non-deterministic ones. Unfortunately, non-deterministic Turing machines (along with any other non-deterministic model of computation) are more of a mathematical tool than anything, and there seems to be no practical way to efficiently solve the problems which would take Non-deterministic Polynomial time (NP-complete problems) as stated by the strong Church-Turing thesis. While it is widely believed that efficiently solving NP-complete problems is impossible, there are some problems which lie in a ‘gray zone’ between NP-COMPLETE and P (i.e. problems which can be solved in deterministic polynomial time): they are believed to be hard for deterministic models (hence, they are not in P), but there is no proof that they are NP-COMPLETE. The most famous of such problems is factorization: with the advent of quantum-computing, which challenges the strong Church-Turing thesis, Shor devised an efficient quantum algorithm for factorizing numbers. While still far from usable in practical cases, its existence proves that one must be extremely careful when talking about the hardness of problems, especially when applied to cryptography, and must always make clear assumptions on the underlying model of computation.

2.2 Finite fields and arithmetic circuits

While computational models typically operate over binary strings, that is, elements of $\{0, 1\}^*$, where $*$ indicates Kleene’s closure, we often want to interpret such strings as elements of some algebraic structure. A *field* is any set equipped with two binary operations, called addition (denoted \oplus) and multiplication (denoted \otimes), which, in simple terms, have all the nice properties of addition and multiplications over real numbers (\mathbb{R} is a field, where $\oplus \equiv +$ and $\otimes \equiv \times$). The most common field used to represent bits is of course the *boolean field* \mathbb{B} , where $\oplus \equiv \text{XOR}$ and $\otimes \equiv \text{AND}$. Bit-strings of length n can be interpreted as elements

of a *finite field* \mathbb{Z}_q , where \oplus and \otimes are defined respectively as integer addition and multiplication modulo q , and $q = p^k$, where p is a prime number and $k \in \mathbb{N}$. Usually, either $p = 2$ and $k = n$, or p is some ‘big’ prime number ($\approx 2^n$) and $k = 1$. When the meaning is clear from the context, we will use $+$ in place of \oplus and omit \otimes , like one would do with real numbers. Sequences of operations over field elements and variables can be neatly represented by *arithmetic circuits*.

Definition 1 (Arithmetic circuit). Given a field \mathbb{F} , some $n, m \in \mathbb{N}$, some constants $a_{1,1}, \dots, a_{m,n} \in \mathbb{F}$, and some variables x_1, \dots, x_n over \mathbb{F} , an *implicit arithmetic circuit* over \mathbb{F} is any formula:

$\phi \equiv c$	with $c \in \mathbb{F}$
$\phi \equiv x$	with x variable over \mathbb{F}
$\phi \equiv \phi_1 \oplus \phi_2$	with ϕ_1 and ϕ_2 arithmetic circuits
$\phi \equiv \phi_1 \otimes \phi_2$	with ϕ_1 and ϕ_2 arithmetic circuits
$\phi \equiv \phi_1^c$	with $c \in \mathbb{F}$ and ϕ_1 arithmetic circuit

An arithmetic circuit which does not contain multiplications and exponentiations by constants is called (*explicit*) *arithmetic circuit*.

Every arithmetic circuit can be represented by a Directed Acyclic Graph (DAG), where the vertices are labeled either with a variable name (*variable vertices*), a constant from the field (*constant vertices*) or one of the operations \oplus (*addition vertices*) and \otimes (*multiplication vertices*, together with the addition vertices are called *operation vertices*). With an analogy to digital circuits, vertices are also called *gates*. Only operation vertices have incoming edges, which must be exactly two, and represent the inputs of the operation, while the outgoing edge will represent the result. It is possible, without affecting the expressive power, to transform an implicit arithmetic circuit into an explicit one by replacing exponentiations (multiplications) by some constant c with a sequence of c multiplications (additions)¹.

Example 1. *Let’s consider the following implicit arithmetic circuit over real numbers:*

$$\phi = x_2(x_1^3 + 4x_2 + 5)$$

We can unroll it into an equivalent (explicit) arithmetic circuit:

$$\hat{\phi} = x_2(x_1x_1x_1 + x_2 + x_2 + x_2 + x_2 + 5)$$

And draw the associated DAG, which is shown in Figure 1.

Any field \mathbb{F} can be extended to an n -dimensional vector space \mathbb{F}^n , for some $n \in \mathbb{N}$, just like one would do with the real numbers. We denote vectors in \mathbb{F}^n

¹On the other hand, such transformation does affect the succinctness of some circuits (unrolling x^c or cx requires $\mathcal{O}(2^c)$ space) and in turn their DAG representation. However, this won’t be a problem for us.




Figure 1: DAG associated to the arithmetic circuit in Example 1.

with lowercase bold letters ($\mathbf{v}, \mathbf{w}, \dots$), and the i th element of a vector \mathbf{v} with \mathbf{v}_i . Vector operations follow their natural definitions depending on the underlying field. We can also introduce matrices over $\mathbb{F}^{n \times m}$ for some $n, m \in \mathbb{N}$, which we denote with bold capital letters ($\mathbf{A}, \mathbf{B}, \dots$). The i th row of a matrix \mathbf{M} is denoted with \mathbf{M}_i , and the j th element of the i th row is denoted with $\mathbf{M}_{i,j}$. Matrix operations also follow their natural definitions over the underlying field. Given $\mathbf{A} \in \mathbb{F}^{n \times m}$, $\mathbf{B} \in \mathbb{F}^{n' \times m'}$, we denote with $(\mathbf{A} \ \mathbf{B})$ their concatenation along the rows, and with $(\mathbf{A}; \mathbf{B}) = (\mathbf{A}^\top \ \mathbf{B}^\top)^\top$ their concatenation along the columns.

A field \mathbb{F} can also be extended to the monovariate polynomial ring $\mathbb{F}[x]$, we will denote polynomials with lowercase letters (p, q, \dots). Operations over polynomials are naturally derived from the underlying field. Vectors and matrices of polynomials are denoted with the usual notation.

Given some $\mathbf{x}, \mathbf{y} \in \mathbb{F}^n$, we can build the unique polynomial:

$$p \mid p \in \mathbb{F}[x] \wedge \deg(p) = n - 1 \wedge \forall i: p(\mathbf{x}_i) = \mathbf{y}_i$$

by using Lagrange interpolation:

$$p = L(\mathbf{x}, \mathbf{y}) = \sum_i \mathbf{y}_i \prod_{j \neq i} \frac{x - \mathbf{x}_j}{\mathbf{x}_i - \mathbf{x}_j}$$

We can extend Lagrange interpolation to any pair of matrices $\mathbf{X}, \mathbf{Y} \in \mathbb{F}^{n \times m}$ by applying L to every row:

$$L(\mathbf{X}, \mathbf{Y}) = (L(\mathbf{X}_1, \mathbf{Y}_1) \dots, L(\mathbf{X}_n, \mathbf{Y}_n))$$

2.3 Rank-1 Constraint systems

Like it happens for boolean formulae and the famous SAT problem, arithmetic circuits can also be seen as a form of constraint whose solution is a set of valid assignments for all the intermediate values in the computation.

Definition 2 (Rank-1 Constraint System). Given a field \mathbb{F} and some $m, n \in \mathbb{N}$, a n/m Rank-1 Constraint System (R1CS) over \mathbb{F} is any triple:

$$\mathcal{C} = (\mathbf{A}, \mathbf{B}, \mathbf{C}) \mid \mathbf{A}, \mathbf{B}, \mathbf{C} \in \mathbb{F}^{n \times m}$$

A *solution* to some R1CS \mathcal{C} is any (column) vector:

$$\mathbf{s} \mid \mathbf{s} \in \mathbb{F}^m \wedge (\mathbf{A}\mathbf{s})(\mathbf{B}\mathbf{s}) = \mathbf{C}\mathbf{s}$$

Any explicit arithmetic circuit with n multiplicative gates and m variables x_1, \dots, x_n can be associated with an $n/(n+m+1)$ R1CS \mathcal{C} which represents the constraints in the circuit, roughly in the following way:

1. Add a new ‘constant variable’ which always assumes value 1.
2. For every multiplicative gate \otimes_i in the circuit, add a new *intermediate* variable t_i (t_n can be denoted y as it represents the circuit output).
3. Define the column vector $\mathbf{x} = (1 \quad x_1 \quad \dots \quad x_m \quad t_1 \quad \dots \quad t_n)^\top$.
4. Express every multiplication gate \otimes_i as an equation in the canonical form:

$$(\mathbf{a}_i \mathbf{x})(\mathbf{b}_i \mathbf{x}) = \mathbf{c}_i \mathbf{x}$$

where $\mathbf{a}_i \mathbf{b}_i \mathbf{c}_i$ will be the i th rows of \mathcal{C}_A , \mathcal{C}_B and \mathcal{C}_C respectively.

Let’s make an example to better understand the process.

Example 2. We have the prime field \mathbb{F}_{13} and the implicit arithmetic circuit:

$$x_2(x_1^3 + 4x_2 + 5)$$

Let’s make the circuit explicit:

$$x_2(x_1x_1x_1 + x_2 + x_2 + x_2 + 5)$$

We can see that there are a total of 3 multiplications in the circuit, and since we have two input variables, our associated R1CS will be a 3/6 R1CS ($2+1+3 = 6$). Let’s explicit all of the intermediate variables:

$$t_1 = x_1x_1 \qquad t_2 = t_1x_1 + 4x_2 + 5 \qquad y = t_2x_2$$

So, our variable vector will be:

$$\mathbf{x} = (1 \quad x_1 \quad x_2 \quad t_1 \quad t_2 \quad y)$$

Now, let’s transform all the equations in canonical form:

$$\begin{aligned} (x_1)(x_1) &= t_1 \\ (t_1)(x_1) + 4x_2 + 5 &= t_2 \iff (t_1)(x_1) = 8 + 9x_2 + t_2 \\ (x_2)(t_2) &= y \end{aligned}$$

Remember that we are working over \mathbb{F}_{13} , so in the second equation, when we bring 4 and 8 to the right side, we have $-4 \equiv 9 \pmod{13}$ and $-5 \equiv 8 \pmod{13}$. We can now extract our R1CS $\mathcal{C} = (\mathbf{A}, \mathbf{B}, \mathbf{C})$:

$$\mathcal{C} = \left(\begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \end{pmatrix}, \begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 \end{pmatrix}, \begin{pmatrix} 0 & 0 & 0 & 1 & 0 & 0 \\ 8 & 0 & 9 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix} \right)$$

By construction, a vector \mathbf{s} is a solution to \mathcal{C} iff every element of \mathbf{s} is assigned to the value derived by fixing x_1, x_2 and following the computation of the original arithmetic circuit. For example, if $x_1 = 2, x_2 = 3$, we have:

$$\begin{aligned} t_1 &= x_1 x_1 = 2 \times 2 = 4 && \equiv 4 \pmod{13} \\ t_2 &= t_1 x_1 + 4x_2 + 5 = 4 \times 2 + 4 \times 3 + 5 = 25 && \equiv 12 \pmod{13} \\ y &= t_2 x_2 = 12 \times 3 = 36 && \equiv 10 \pmod{13} \end{aligned}$$

Therefore our solution vector will be:

$$\mathbf{s} = (1 \quad 2 \quad 3 \quad 4 \quad 12 \quad 10)$$

It is a bit tedious, but easy, to verify that indeed $(\mathbf{A}\mathbf{s})(\mathbf{B}\mathbf{s}) = \mathbf{C}\mathbf{s}$.

2.4 Quadratic Arithmetic Programs

A problem with R1CS is that solutions have size linear in the number of multiplication gates of the corresponding arithmetic circuit. This can be solved by using Quadratic Arithmetic Programs.

Definition 3 (Quadratic Arithmetic Program). Given a field \mathbb{F} and some $n, m \in \mathbb{N}$, a n/m Quadratic Arithmetic Program (QAP) over \mathbb{F} is any quadruple:

$$\mathcal{Q} = (t, \mathbf{v}, \mathbf{w}, \mathbf{y}) \mid t \in \mathbb{F}[x] \wedge \mathbf{v}, \mathbf{w}, \mathbf{y} \in \mathbb{F}[x]^n$$

For which it holds that:

$$\forall i: \deg(\mathbf{v}_i) + 1 = \deg(\mathbf{w}_i) + 1 = \deg(\mathbf{y}_i) + 1 = \deg(t) = m$$

A valid assignment to a QAP \mathcal{Q} is any vector:

$$\mathbf{s} \in \mathbb{F}^n \mid (\mathbf{v}\mathbf{s})(\mathbf{w}\mathbf{s}) - \mathbf{y}\mathbf{s} \bmod t = 0$$

Then, the polynomials $p = (\mathbf{v}\mathbf{s})(\mathbf{w}\mathbf{s}) - \mathbf{y}\mathbf{s}$ and $h = \frac{p}{t}$ are a solution to \mathcal{Q} .

Just like it was possible to represent any n/m arithmetic circuit ϕ with an $n/(n+m+1)$ R1CS \mathcal{C} , we can, in turn, represent any n/m R1CS $\mathcal{C} = (\mathbf{A}, \mathbf{B}, \mathbf{C})$ with a n/m QAP \mathcal{Q} . First, we choose some arbitrary $\mathbf{z} \in \mathbb{F}^n \mid \forall i, j: \mathbf{z}_i \neq \mathbf{z}_j$ (usually, $\mathbf{z} = (1 \quad \dots \quad n)$). Let $\mathbf{Z} \in \mathbb{F}^{m \times n} \mid \forall i: \mathbf{Z}_i = \mathbf{z}$, then:

$$\mathcal{Q} = (t, \mathbf{v}, \mathbf{w}, \mathbf{y}) = \left(\prod_i (x - \mathbf{z}_i), L(\mathbf{Z}, \mathbf{A}^\top), L(\mathbf{Z}, \mathbf{B}^\top), L(\mathbf{Z}, \mathbf{C}^\top) \right)$$

To make things more clear, let's make an example:

Example 3. We want to compute the 3/6 QAP $\mathcal{Q} = (t, \mathbf{v}, \mathbf{w}, \mathbf{y})$ associated with the 3/6 R1CS $\mathcal{C} = (\mathbf{A}, \mathbf{B}, \mathbf{C})$ that we derived in Example 2. First, we set:

$$\mathbf{z} = (1 \quad 2 \quad 3) \quad \mathbf{Z} = (\mathbf{z}; \mathbf{z}; \mathbf{z}; \mathbf{z}; \mathbf{z}; \mathbf{z})$$

Then, we compute the target polynomial t , the left and right input constraint polynomial vectors \mathbf{v} and \mathbf{w} , and the output constraint polynomial vector \mathbf{y} (remember, we are working over \mathbb{F}_{13}). Notice how the 2nd, 4th and 5th columns of \mathbf{A} form the canonical basis of \mathbb{F}_{13}^3 , and since L is a linear operator, we can express all other polynomials as linear combinations of $L(\mathbf{z}, \mathbf{A}_2^\top)$, $L(\mathbf{z}, \mathbf{A}_4^\top)$ and $L(\mathbf{z}, \mathbf{A}_5^\top)$:

$$t = (x-1)(x-2)(x-3) = (x+12)(x+11)(x+10) = x^3 + 7x^2 + 11x + 7$$

$$\mathbf{v} = L(\mathbf{Z}, \mathbf{A}^\top) = (L(\mathbf{z}, \mathbf{A}_1^\top) \quad \cdots \quad L(\mathbf{z}, \mathbf{A}_6^\top)) = \begin{pmatrix} 0 \\ 7x^2 + 4x + 3 \\ 0 \\ 12x^2 + 4x + 10 \\ 7x^2 + 5x + 1 \\ 0 \end{pmatrix}^\top$$

$$\mathbf{w} = (0 \quad \mathbf{v}_2 + \mathbf{v}_4 \quad \mathbf{v}_5 \quad 0 \quad 0 \quad 0) = \begin{pmatrix} 0 \\ 6x^2 + 8x \\ 7x^2 + 5x + 1 \\ 0 \\ 0 \\ 0 \end{pmatrix}^\top$$

$$\mathbf{y} = (8\mathbf{v}_4 \quad 0 \quad 9\mathbf{v}_4 \quad \mathbf{v}_2 \quad \mathbf{v}_4 \quad \mathbf{v}_5) = \begin{pmatrix} 5x^2 + 6x + 2 \\ 0 \\ 4x^2 + 10x + 12 \\ 7x^2 + 4x + 3 \\ 12x^2 + 4x + 10 \\ 7x^2 + 5x + 1 \end{pmatrix}^\top$$

Recall that a possible solution to the R1CS was:

$$\mathbf{s} = (1 \quad 2 \quad 3 \quad 4 \quad 12 \quad 10)$$

Let's check if it is also a valid assignment for the QAP:

$$\begin{aligned} p &= (\mathbf{v}\mathbf{s})(\mathbf{w}\mathbf{s}) - \mathbf{y}\mathbf{s} = (3x^2 + 6x + 6)(7x^2 + 5x + 3) - (12x^2 + 7x + 11) \\ &= 8x^4 + 5x^3 + 4x^2 + 2x + 7 \end{aligned}$$

$$h = \frac{p}{t} = \frac{8x^4 + 5x^3 + 4x^2 + 2x + 7}{x^3 + 7x^2 + 11x + 7} = 8x - 51 + \frac{273x^2 + 507x + 364}{x^3 + 7x^2 + 11x + 7} = 8x + 1$$

Since:

$$ht = (8x + 1)(x^3 + 7x^2 + 11x + 7) = 8x^4 + 5x^3 + 4x^2 + 2x + 7 = p$$

this means that p and h are a solution to the QAP, and \mathbf{s} is a valid assignment.

One might wonder how a solution (p, h) to a QAP is more succinct than the corresponding valid assignment \mathbf{s} of the associated R1CS: as a matter of fact, given an n/m arithmetic circuit, \mathbf{s} has size $n + m + 1$, while p can have degree (and therefore encoding size) $2(n-1)$. Furthermore, in a typical circuit, $n \gg m$, so (p, h) would approximately be twice the size of \mathbf{s} when encoded. Now, $p = ht \implies \forall x: p(x) = h(x)t(x)$; if we are working over a big field (say, $|\mathbb{F}| \approx 2^{256}$), it is hard to find even a single value of x for which the equation holds. This means that we can accept as a solution, with high confidence (although not certainty) any couple of values x, y such that $y \bmod t(x) = 0$.

Summing up: if $y \bmod t(x) = 0$, we are *almost sure* that y has been derived by computing $p(x)$, where p is a solution to our QAP. But if p is a solution to the QAP, then it derives from a valid assignment \mathbf{s} to the associated R1CS, which in turn derives from a valid computation of the original arithmetic circuit.

2.5 Hash functions

Hash functions are a fundamental tool in many fields of computer science, and cryptography is arguably the most prominent. Formally, an hash function is any function $H: \{0, 1\}^* \mapsto \{0, 1\}^n$, that is any function which maps arbitrarily long *messages* to fixed-size *digests*. From the definition, it is immediate to see that there are an infinite number of messages which map to the same digest. While an operation like truncation is a (very simple) hash function, in cryptography we are interested in functions that provide additional guarantees: the assumption is that a digest should represent a message in a one-way fashion: while there are infinite messages which map to the same digest, it must be hard to find them. Ideally, a cryptographic hash function should behave like a perfect random function. This is of course impossible, as the output of an hash function must only depend deterministically on its input; the aim then is to build functions which are hard to distinguish from a random function.

Definition 4 (Cryptographic hash function). Given $n \in \mathbb{N}$, an n (-bit) *cryptographic hash function* (CHF) is any function $H: \{0, 1\}^* \mapsto \{0, 1\}^n$ which satisfies the following properties:

- **Collision resistance:** It is hard to find two messages m_1, m_2 such that $H(m_1) = H(m_2)$.
- **Preimage resistance:** Given some digest h , it is hard to find a message m such that $H(m) = h$ (H is a one-way function).
- **Second preimage resistance:** Given some message m_1 , it is hard to find a message m_2 such that $H(m_1) = H(m_2)$.

While some of the requirements might seem redundant (for example, if it is hard for an attacker to find a collision for chosen messages, it must be hard when one is fixed), the difference usually lies in how exactly we define hardness for each property. For collision resistance, an ideal CHF requires about $2^{n/2}$ evaluations to find a collision (birthday paradox), while for preimage resistance it would

require about 2^n evaluations. Typically, a CHF is built by applying some known secure constructions to functions which are simpler to devise.

Definition 5 (Pseudorandom keyed permutation). Given $l, n \in \mathbb{N}$, an l/n -bit pseudorandom keyed permutation (PKP) is any bijective function:

$$F: \{0, 1\}^l \times \{0, 1\}^n \mapsto \{0, 1\}^l$$

which is hard to distinguish from an uniform random distribution.

PKPs are often built by iterating a keyed permutation F for some number r of rounds, since F by itself might be relatively easy to invert. A block cipher is a pseudorandom keyed permutation which changes the key being used in each round through a key-scheduling function. Unkeyed permutations can be derived from keyed ones simply by fixing the key to some arbitrary value.

Definition 6 (One-way compression function). Given $l_1, l_2, n \in \mathbb{N}$, an $l_1/n/l_2$ -bit one-way compression function (OWCF) is any function:

$$F: \{0, 1\}^{l_1} \times \{0, 1\}^n \mapsto \{0, 1\}^{l_2}$$

There are many known ways to build OWCFs from pseudorandom keyed permutations, and, in turn, CHFs from OWCFs. We will introduce the Davies-Meyer and the Merkle-Damgård constructions respectively, as those are the ones of interest to us.

Theorem 1 (Davies-Meyer construction). Given a l/n pseudorandom keyed permutation E , some $i, k \in \mathbb{N}$, some $v \in \{0, 1\}^l$, and some $m \in \{0, 1\}^{kn}$, then any function F_E such that:

$$F_{E,i}(v, m) = \begin{cases} E(v, m_{1\dots n}) & i = 1 \\ E(F_{E,i-1}(v, m), m_{i(n-1)\dots in}) & 2 \leq i \leq k \end{cases}$$

$$F_E = F_{E,k}$$

is a $l/kn/l$ OWCF.

Theorem 2 (Merkle-Damgård construction). Given a $l_1/n/l_2$ OWCF F , some $k \in \mathbb{N}$, some $v \in \{0, 1\}^{l_1}$, some $m \in \{0, 1\}^*$ and some padding function:

$$P(m): \{0, 1\}^{|m|} \mapsto \{0, 1\}^{|m| + (-|m| \bmod n) + kn}$$

such that, $\forall m, m' \in \{0, 1\}^*$:

$$(|m| = |m'| \Rightarrow |P(m)| = |P(m')|) \wedge (|m| \neq |m'| \Rightarrow m_{|P(m)|} \neq m'_{|P(m')|})$$

then any function H_F such that:

$$H_{F,i}(v, m) = \begin{cases} F(v, m_1) & i = 1 \\ F(H_{F,i-1}(v, m), m_i) & 1 < i \leq |P(m)| \end{cases}$$

$$H_F = H_{F,|P(m)|}$$

is a cryptographic hash function.

2.6 Tree hash modes

An important application of CHF is in *prover-verifier games*: for any message m , the digest $h = H(m)$, where H is an n CHF, can be used as a *binding commitment* for m : a verifier is convinced that the prover knows m simply by asking him to share h , with overwhelming confidence ($\approx 1 - \frac{1}{2^n}$). While often referred to as if they were humans, provers and verifiers are formally described by some model of computation, usually deterministic Turing machines, which often can only harness a limited amount of resources (time and space), typically at most polynomial in the size of the game instance statement².

If the prover wants to commit to a list of k messages, a possibility would be to share with the verifier the hash of every message: this would require a $\mathcal{O}(k)$ communication cost and a $\mathcal{O}(k)$ verification cost. A slightly better alternative would be for the prover to share $H(\{m_1, \dots, m_k\})$: the communication cost would only be $\mathcal{O}(1)$, but verification would still cost $\mathcal{O}(k)$.

Definition 7 (Merkle Tree). Given some $k \in \mathbb{N}$, a CHF H and a set of messages $S = \{m_1, \dots, m_{s^{k-1}} \mid \forall i: m_i \in \{0, 1\}^*\}$, a *Merkle Tree (MT)* is a complete binary tree of height k such that:

1. The leaf nodes $\nu_1, \dots, \nu_{2^{k-1}}$ contain $H(m_1), \dots, H(m_{2^{k-1}})$.
2. Every other node ν contains $H(\nu_l, \nu_r)$, where ν_l is the left child of ν and ν_r is the right child of ν .

By using Merkle trees, the prover only needs to send to the verifier, as a commitment for some message m_i among $k = 2^{\lceil \log_2(k) \rceil}$ messages, the contents of the co-path from the leaf containing m_i to the root (plus the hash of m_i): this requires just $\mathcal{O}(\log_2(k))$ communication effort and $\mathcal{O}(\log_2(k))$ verification effort. Another advantage of Merkle trees is that bottom-up construction is very easy to parallelize, and its usefulness can be appreciated even more when considering a scenario where different messages actually belong to different provers.

Definition 8 (Augmented Binary tRee). Given some $k \in \mathbb{N}$, a CHF H , and a set of messages $S = \{m_1, \dots, m_{2^{k-1}+2^{k-2}-1} \mid \forall i: m_i \in \{0, 1\}^*\}$, an *Augmented Binary tRee (ABR)* is a complete binary tree of height k augmented with *middle* nodes, such that:

1. The leaf nodes $\nu_1, \dots, \nu_{2^{k-1}}$ contain $H(m_1), \dots, H(m_{2^{k-1}})$.
2. There are no middle nodes in the leaf layer.
3. The middle nodes $\nu_{2^{k-1}+1}, \dots, \nu_{|S|}$ contain $H(m_{2^{k-1}+1}), \dots, H(m_{|S|})$.
4. Every other node ν contains $H(\nu_l \oplus \nu_m, \nu_r \oplus \nu_m) \oplus \nu_r$, where ν_l is the left child of ν , ν_r is the right child of ν , and ν_m is the middle child of ν , or 0 if ν doesn't have a middle child.

²Although humans can be assimilated to a computational model, it is not easy to formalize the eventuality of the prover threatening the verifier to make him accept his proof..

Notice the use of the \oplus operation inside the ABR: while messages of length n are usually treated as elements of $\{0, 1\}^n$, they can also be treated as n -bit integers over some field \mathbb{F}_q : if $q = 2^n$, then \oplus means bitwise XOR (i.e. addition in \mathbb{F}_2), and if $q = p$ for some prime p , then \oplus means addition in the field \mathbb{F}_p .

ABRs can store 50% more messages than Merkle Trees for the same height, resulting in the same number of calls to H , at the cost of performing 3 additional \oplus operations for every call (usually, $TIME(\oplus) \ll TIME(H)$).

3 ZK-SNARK

We saw in Section 2 how a prover can convince a verifier about the knowledge of some message m , with a high confidence and a small communication effort, by using a CHF H . However, the underlying assumption was that m is known by the verifier: when the prover sends h , the verifier can check whether $H(m) = h$ and therefore accept or reject. In this Section, we will see how a prover can convince a verifier without the need to disclose possibly secret information. In particular, we will focus on provable computation, that is, when the prover wants to convince the verifier that he correctly computed some function.

3.1 Zero Knowledge Proofs

Before diving into provable computation, we must introduce the more general concept of Zero Knowledge Proof system.

Definition 9 (Zero-Knowledge Proof). Given a prover P and a verifier V , a secret x , known only to P , and some statement σ of whose truth P wants to convince V by means of some proof π , we call a Zero-Knowledge Proof (ZKP) system any protocol which satisfies the following properties:

- **Soundness:** $\neg\sigma \implies V(\pi) = \perp$.
- **Completeness:** $\sigma \implies V(\pi) = \top$.
- **Zero-Knowledge:** It is *hard* for V to derive x given σ and π .

While formal proofs have been known for millenia, only in the last century, with the advent of modern cryptography, researchers started considering the possibility of having proofs of statements which, while able to convince someone of their truth, didn't leak information about how they were obtained. Zero-Knowledge systems prove particularly useful in *Argument of Knowledge* scenarios (ZK-ARK): the prover P wants to convince the verifier V that he knows a solution to some computational problem, assuming there is one, without revealing it. This becomes particularly handy when dealing with hard problems, like committing to some message m by computing some secure hash function $h = H(m)$. Since H is a OWF, it would be too hard for V to 'make-up' some fake m' which hashes to h . It must be noted though that known ZK-ARK systems do not guarantee the formal soundness of the proof: there is a small probability that, given some false statement σ and an (invalid) proof π , then $V(\pi) = \top$, so it is important to keep this probability small (say, 2^{-128}). There are other nice additional properties that zero-knowledge systems might satisfy, making them even more interesting.

Definition 10 (ZK-SNARK). Given a prover P , a verifier V , a statement σ , and a ZK-ARK system to produce a proof π , if the system is:

- **Succinct:** $SPACE(\pi) = o(\log(\sigma))$.

- **Non-interactive:** The only communication required by the system is the exchange of σ and π .

then it is a Zero-Knowledge Non-interactive ARgument of Knowledge (ZK-SNARK) system.

Succintness is an important property in a blockchain scenario, since we cannot afford to use too much space to store the proofs, and non-interactivity of the process allows for efficient verification when multiple parties are involved.

One of the most important applications of ZK-SNARK systems is in *provable computation*, where the prover wants to convince the verifier that he correctly performed some computation (e.g. a cryptocurrency transaction). A very famous ZK-SNARK system for verifiable computation is the *Pinocchio* protocol, which was the first one efficient enough to be practical.

3.2 The Pinocchio Protocol

Pinocchio is composed of many different components, and requires quite a bit of mathematical background to be fully understood. We will not go into all of the mathematical and cryptographic details of the protocol, especially the ones involving *elliptic curves*, but we will still try to give a good idea of how the protocol works, and ultimately what determines its computational complexity.

Pinocchio does not allow the encoding of arbitrary languages, i.e. it is not Turing complete, but we are restricted to arithmetic circuits over an arbitrary prime field \mathbb{F}_p . The main limitation arising from this restriction is that we cannot express unbounded computation (e.g. loops whose exit condition depends on some non-constant value) in this framework. This issue can be mitigated by writing a circuit compiler in a Turing complete language which is able to synthesize parametrized arithmetic circuits on the fly.

The reason we translate arithmetic circuits into R1CS is that they explicit all of the computation in terms of linear combinations, which allows us to use Lagrange interpolation to build polynomials over them.

With all this in mind, the flow of the Pinocchio protocol is as follows:

1. Encode an algorithm as an arithmetic circuit over some prime field \mathbb{F}_p .
2. Compute the associated R1CS.
3. Compute the associated QAP.
4. Generate random values and instantiate an homomorphic encryption mapping (using elliptic curves) which depends on those values.
5. Generate the proof by encrypting the QAP with the homomorphic mapping.
6. Map again the proof to a new homomorphic space, and finally verify it.

Due to the homomorphism of the mappings and the properties of QAPs and R1CSs, if the verification is successful, it means that the original algorithm was in fact correctly executed, with high probability. If the verification fails, then certainly the original algorithm was not executed correctly. The verifier learns cannot learn any additional information from this process without performing an infeasible amount of work, therefore this is indeed a ZK-SNARK protocol.

- 4 The libsnark library
- 5 Experiments
- 6 Conclusions and future directions